

Predicting Medical Insurance Charges using Linear Regression and Random Forest

BY GRISHMA VEMIREDDY



Introduction

Problem statement:

Predict a person's medical insurance charges based on their demographic data and lifestyle.

Purpose:

It supports risk assessment and customer profiling.

Helps understand how insurance companies estimate premiums.



```
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

Dataset Overview

Source: Kaggle ([Medical Cost Personal Datasets](#))

Data description:

1,338 total observation

No missing values

7 features:

age, sex, bmi, children, smoker, region, charges
(target variable).

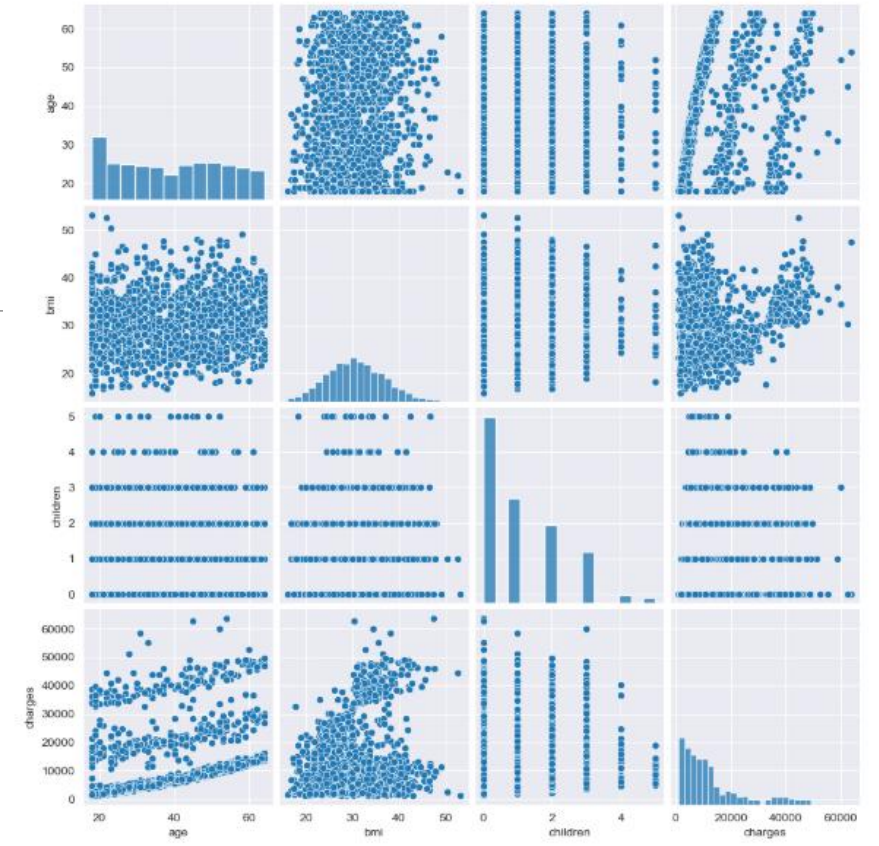


Methodology

First performed data cleaning and data exploration, scaled the data, then split data to train and test models.

```
#conver categorical to numerical
df_encoded = data_df.copy()
df_encoded['sex'] = df_encoded['sex'].map({'male': 0, 'female': 1})
df_encoded['smoker'] = df_encoded['smoker'].map({'no': 0, 'yes': 1})
df_encoded['region'] = df_encoded['region'].map({'southwest': 0, 'southeast': 1})
```

```
sns.heatmap(df_encoded.corr(), annot=True)
plt.show()
```



```
# scale data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

#Splt to train and test 80-20 split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2)
```

Methodology

Two regression models:

1. Linear regression

Implemented a basic model to predict charges.

2. Random Forest

I initially passed in parameters of

n_estimators=100

max_depth=10

Then tuned and tested the model with different parameter values to see any performance improvement

n_estimators=300

max_depth=20

```
#Linear regression
lm = LinearRegression()
lm.fit(X_train,y_train)
```

▼ LinearRegression

LinearRegression()

```
# Random Forest
rfc = RandomForestRegressor(n_estimators=100,max_depth=10,random_state=42)
rfc.fit(X_train, y_train)
```


▼ RandomForestRegressor

RandomForestRegressor(max_depth=10, random_state=42)



Results and Visualization


Linear Regression



```
# RMSE and R^2
print('RMSE:', np.sqrt(mean_squared_error(y_test, predictions)))
print('R2 Score: ', r2_score(y_test, predictions))
```

RMSE: 5799.5870914383595
R2 Score: 0.7833463107364536


Random Forest



```
print('RMSE:', np.sqrt(mean_squared_error(y_test, rfc_pred)))
print('R2 Score: ', r2_score(y_test, rfc_pred))
```

RMSE: 4572.110965856546
R2 Score: 0.8653502770474271





Conclusion

Overall, the random forest regression model provided with the best predictions performance with high r-squared score and lower RMSE compare with the linear regression model.

While the Linear Regression model gives a good baseline, the Random Forest model is much more effective at capturing complex relationships within the data.



Thank You!

